

Intelligent Drug Sales Forecasting and Healthcare -Insights with Artificial Intelligence

Final Year Research Proposal

**B R G Lakmal
CS/2020/015**

Supervisor: Prof N G J Dias

1. Project Charter

Pharmaceutical sales forecasting plays a critical role in ensuring that medications are available when and where they are needed, helping to prevent shortages and improve patient outcomes. Traditional forecasting methods often fail to account for complex patterns caused by seasonal demand, public health events, environmental factors, and human behavior. Moreover, most existing models provide predictions as numbers without offering explanations or actionable insights, making it difficult for stakeholders to understand the underlying reasons behind fluctuations in demand.

This project aims to create an AI-driven drug sales forecasting and advisory system that not only predicts future drug demand but also provides context-aware, explainable insights and recommendations. By leveraging time-series analysis, machine learning, deep learning, and fine-tuned language models (LLMs), the system will identify trends, anticipate potential challenges, and provide guidance for proactive decision-making. The model will be capable of analyzing historical sales data, health-related literature, environmental factors, and other contextual information to produce both quantitative forecasts and qualitative insights.

The main stakeholders for this project are pharmacists, healthcare providers, pharmaceutical companies, public health authorities, and patients who rely on timely and accurate medication availability. By providing actionable insights, the system aims to:

- Enhance inventory management and reduce the risk of stockouts or overstocking.
- Offer explanations for changes in demand, helping stakeholders understand the influence of seasonal patterns, health events, and environmental factors.
- Support public health monitoring by identifying potential health risks associated with changing drug demand.
- Improve decision-making across healthcare and pharmaceutical sectors through data-driven recommendations.

To achieve this, the project assumes that historical sales data and health-related Knowledge sources are representative of real-world trends. It also assumes access to sufficient computational resources for training advanced models and fine-tuning the LLM. Furthermore, it assumes that environmental and epidemiological data can be reliably integrated into the system to improve predictive accuracy.

Key challenges to consider include:

- **Data variability and quality:** Sales data may be incomplete, inconsistent, or collected from multiple heterogeneous sources.
- **Integration of external factors:** Combining environmental, epidemiological, and health literature data into the forecasting system may be complex.
- **Model interpretability:** Translating advanced predictions into clear, actionable insights for non-technical users.
- **Privacy and ethical compliance:** Ensuring that sensitive sales and health data are securely handled in accordance with regulations.
- **Adaptability and scalability:** Developing a system that can handle changes in drug categories, regions, and emerging health trends without losing accuracy.

Successfully addressing these challenges will result in a robust, intelligent, and explainable drug forecasting and advisory system that supports both operational efficiency in pharmaceutical supply chains and public health decision-making. Ultimately, the project seeks to improve medication availability, patient care, and proactive management of health risks in a reliable and sustainable way.

2. Background

Accurate forecasting of pharmaceutical sales is critical for ensuring that medications are available when and where they are needed, preventing shortages, and supporting public health. Traditional sales forecasting methods, such as simple time-series models or historical averaging, often fail to capture the complex patterns influenced by seasonal demand, public health events, environmental factors, and human behavior. These limitations can lead to overstocking, shortages, or misallocation of resources, which in turn can affect patient care and healthcare outcomes.

Recent advancements in artificial intelligence (AI) and machine learning (ML) have shown great potential in improving pharmaceutical sales predictions. Techniques such as random forests, gradient boosting and Machine Learning models have been applied to capture non-linear and long-range dependencies in sales data, allowing for more accurate forecasts. These models, however, often act as “black boxes,” providing predictions without explanations, which can limit trust and practical usability for pharmacists, healthcare providers, and decision-makers.

To address this, the integration of fine-tuned language models (LLMs) and agentic AI has been proposed as a solution. Fine-tuning LLMs on health-related literature, drug manuals, and historical sales patterns enables the system to provide interpretable insights and actionable recommendations. For example, the model can explain why a particular drug category may see higher demand, identify potential health risks, and suggest preventive measures for pharmacists and healthcare providers.

Several studies have demonstrated the effectiveness of AI in sales and demand forecasting. For instance, research applying LSTM models to pharmaceutical sales data achieved high accuracy in predicting weekly drug demand across multiple categories.

Another study incorporated environmental and contextual factors, such as temperature, epidemic alerts, and public events, into ML-based forecasts, improving prediction reliability by over 15% compared to baseline methods.

Despite these advancements, challenges remain in data variability, model generalization, and Pipeline integration, and integration of external contextual factors. Sales data often comes from multiple sources with varying quality, granularity, and completeness.

Additionally, environmental and epidemiological factors must be properly aligned and weighed up to ensure accurate forecasting. Finally, translating complex model outputs into actionable insights that stakeholders can trust requires careful design and the integration of explainable AI techniques.

By combining time-series forecasting, machine learning, deep learning, and fine-tuned LLMs, this project aims to create a robust, explainable, and context-aware drug sales forecasting and advisory system and Auto CI/CD Pipeline integration. Such a system can not only predict future demand but also provide insights into why trends are occurring, what health or environmental risks may be associated, and what actions should be taken, supporting better decision-making for pharmacists, healthcare providers, pharmaceutical companies, and public health authorities.

3. Aims

The aim of this project is to develop an AI-driven Autonomous drug sales forecasting and advisory system that not only predicts future medication demand accurately but also provides interpretable, actionable insights for stakeholders. The system will leverage time-series models, machine learning, deep learning, Autonomous Pipeline and fine-tuned language models (LLMs) to forecast sales for different drug categories and explain the factors influencing demand.

By integrating contextual information such as health trends, seasonal patterns, environmental conditions, and historical sales data, the system will help pharmacists, healthcare providers, and pharmaceutical companies understand why certain drugs are likely to experience increased or decreased demand. Additionally, it will provide recommendations for inventory management, patient advisories, and public health interventions, enabling proactive decision-making.

The overall goal is to create a robust, scalable, and explainable solution that improves drug availability, minimizes stockouts, mitigates potential health risks, and supports better resource planning in healthcare and pharmaceutical settings. This system is designed to be usable across different regions and organizations, including those with limited access to advanced analytics tools.

4. Objectives

- **Collect and consolidate reliable data sources:** Gather historical pharmaceutical sales data according to ATC classification Categories, health-related literature, environmental data (e.g., weather, epidemics), and contextual information relevant to drug demand.
- **Preprocess and clean datasets:** Handle missing values, normalize data, create lagged and engineered features, and align contextual inputs to prepare the data for modeling.
- **Develop and train predictive models:** Implement classical time-series models (ARIMA, SARIMAX), machine learning models (Random Forest, XGBoost, Gradient Boosting), and deep learning models (LSTM, GRU, TCN, Transformers) to forecast weekly sales for different drug categories.
- **Fine-tune domain-specific language models (LLMs):** Train LLMs on health-related literature, pharmaceutical guidelines, and historical sales context to provide explainable insights and recommendations.
- **Integrate agentic AI for adaptive forecasting:** Combine outputs from classical, ML, and DL models with the fine-tuned LLM and MLOps Automation to produce context-aware predictions and select the most reliable forecast automatically.

- **Validate model performance:** Evaluate predictive accuracy using metrics such as MAE, RMSE, and MAPE, and assess the reliability of the generated explanations and recommendations.
- **Generate actionable advisory reports:** Produce interpretable reports indicating predicted sales trends, reasons behind demand changes, potential health risks, environmental influences, and suggested preventive or corrective actions.
- **Develop a prototype application:** Build a user-friendly web-based interface to allow stakeholders to explore forecasts, receive automated recommendations, and interact with the advisory system in real-world settings.

5. Feasibility Analysis

- **Data Availability:**
Historical pharmaceutical sales data, ATC drug classification datasets, health-related literature, and environmental/contextual data are accessible from multiple sources. Publicly available drug sales records, pharmacy chains, health ministry reports, and research publications provide a solid foundation for training and testing the forecasting and advisory models. This diverse data ensures that the system can generalize across multiple drug categories, seasonal patterns, and contextual influences.
- **Technological Feasibility:**
The project leverages well-established techniques in time-series forecasting, machine learning, and deep learning (e.g., ARIMA, SARIMAX, Random Forest, XGBoost, LSTM, GRU, TCN, Transformers). Additionally, fine-tuned language models (LLMs) and agentic AI frameworks will be integrated to provide explainable insights and adaptive predictions. Popular platforms such as TensorFlow, PyTorch, Scikit-learn, and Hugging Face Transformers are widely supported and provide robust tools to implement these models efficiently.
- **Computational Resources:**
The computational requirements for training ML/DL models and fine-tuning LLMs can be met using cloud-based services such as Google Colab, AWS, or Azure, which offer scalable processing power and GPU acceleration. This allows efficient model training without the need for costly local hardware. Additionally, modern desktop or laptop GPUs (e.g., NVIDIA 1660 Super, RTX 4060) can be leveraged for preliminary experiments and model fine-tuning, making the project feasible for both research and practical deployment.

- **Operational Feasibility:**

The system can be deployed as a web-based application, providing pharmacists, healthcare providers, and public health authorities with an intuitive interface for exploring predictions, receiving recommendations, and monitoring trends in real-time. The integration of explainable AI ensures that outputs are interpretable, supporting informed decision-making in operational settings.

- **Economic Feasibility:**

Utilizing cloud resources and open-source frameworks minimizes the cost of development. Data acquisition can largely rely on publicly available or partnership-based sources, reducing expenses. The project's potential to improve inventory management, reduce stockouts, and support public health interventions may also provide significant cost savings for healthcare providers and pharmaceutical companies.

6. Methodology

Data Collection and Preprocessing:

The first step will be to gather historical pharmaceutical sales data, including weekly sales by ATC drug categories, as well as health-related literature, environmental data (e.g., weather patterns, epidemics), and contextual information. The data will be cleaned to handle missing values, remove outliers, and standardize formats. Feature engineering will include generating lagged features, moving averages, seasonal indicators, and relevant contextual variables. Data will be normalized to improve model performance, and train-test splits will be applied to ensure unbiased evaluation.

Model Selection and Training:

The project will leverage a hybrid modeling approach:

- **Classical Time-Series Models:** ARIMA and SARIMAX will capture basic trends and seasonality, providing a reliable baseline.
- **Machine Learning Models:** Random Forest, XGBoost, and Gradient Boosting will model non-linear relationships between features, capturing complex patterns missed by classical models.
- **Deep Learning Models:** LSTM, GRU, Temporal Convolutional Networks (TCN), and Transformers will be applied to capture sequential dependencies and long-term patterns in sales data.

Each model will be trained and validated using historical data, and hyperparameter tuning will be performed to optimize performance. Pre-trained models will be fine-tuned if applicable, leveraging transfer learning for enhanced accuracy.

Applying MLOPS Pipeline Integration in order to Update and Train entire Model Flow with new available data.

Fine-Tuned Language Model (LLM) Integration:

A domain-specific LLM will be fine-tuned on pharmaceutical literature, health guidelines, and historical sales trends. The LLM will provide explainable insights for predictions, including:

- Reasons why a particular drug category is likely to see increased demand.
- Potential health or environmental factors influencing sales.
- Preventive measures or recommendations for pharmacists and healthcare providers.

Agentic AI Integration:

An agentic AI system will orchestrate all predictive models, deciding which model's output is most reliable for a given scenario. The agent will:

- Select classical, ML, or DL models depending on forecast horizon and data conditions.
- Consult the fine-tuned LLM for explanations and advisory reports.
- Provide actionable recommendations to end-users.

Model Evaluation:

The performance of predictive models will be assessed using metrics such as MAE, RMSE, MAPE, and R^2 , while the quality of LLM-generated explanations will be evaluated for clarity, relevance, and usefulness. Cross-validation and back testing will ensure that models generalize well to unseen data.

Deployment and Testing:

The final system will be deployed as a user-friendly web application, allowing stakeholders to:

- Input drug category and date to receive sales forecasts.
- Access explanatory insights and advisory reports generated by the fine-tuned LLM.
- Compare outputs from classical, ML, and DL models.

Validation and Feedback (based on availability):

The system's accuracy, interpretability, and usability will be validated by gathering feedback from pharmacists, healthcare providers, and public health professionals. This will ensure that forecasts are actionable, and explanations are clear and trustworthy. Continuous feedback will inform model updates and improve system performance over time.

7. Time Frame

Research is planned to be completed within 36 weeks. The key milestones are outlined as follows:

Task	Week
Selecting Topic C Defining Scope	1-2
Literature Review C Background Study	3-4
Proposal Preparation	1-4
Proposal Presentation	5-6
Data Collection C Preprocessing	7-10
Model Selection and Training (ARIMA, ML, DL)	11-12
LLM Fine-Tuning and Explainability Integration	13-16
Agentic AI Integration	17-18
Model Evaluation and Testing	19-21
Application Deployment and Testing	22-24
Final Adjustments C Feedback Incorporation	24-32

8. References

- [1] S. Abbas, A. Qaisar, M. S. Farooq, M. S. M. Ahmad, and M. A. Khan, “Smart Vision Transparency: Efficient Ocular Disease Prediction Model Using Explainable Artificial Intelligence,” *Sensors*, vol. 24, no. 20, 2024.
- [2] M. S. Khan, N. Tafshir, K. N. Alam, A. R. Dhruva, M. M. Khan, A. A. Albraikan, and F. A. Almalki, “Deep Learning for Ocular Disease Recognition: An Inner-Class Balance,” *Computational Intelligence and Neuroscience*, 2022.
- [3] T. Nazir *et al.*, “Detection of Diabetic Eye Disease from Retinal Images Using a Deep Learning-Based CenterNet Model,” *Sensors*, vol. 21, no. 16, 2021.
- [4] T. Gamage *et al.*, “Age-Related Retinal Disease Identification Using Image Processing and Deep Learning,” in *Proc. 4th Int. Conf. Computing and Artificial Intelligence (CCAI)*, 2024.
- [5] K. H. Wijesinghe *et al.*, “Identification of Diabetic Related Eye Diseases Using Deep Learning,” in *Proc. 5th Int. Conf. Advancements in Computing (ICAC)*, 2023.
- [6] S. B. Ekanayake, G. Lakmal, A. Perera, M. Nasmeen, P. Vimanshani, and M. D. Chandrasena Atampalage, “Predicting Medical Drug Sales in a Specific Area for Categorical Drugs Using Time Series Forecasting,” in *Proc. Int. Research Conf. Smart Computing and Systems Engineering (SCSE)*, Colombo, Sri Lanka, Apr. 2025, pp. 1–6, doi:10.1109/SCSE65633.2025.11031056.
- [8] P. Hewamalage, C. Bergmeir, and K. Bandara, “Recurrent Neural Networks for Time Series Forecasting: Current Status and Future Directions,” *International Journal of Forecasting*, vol. 37, no. 1, pp. 388–427, 2021.
- [9] M. Lim and A. Zohdy, “Explainable Artificial Intelligence for Business Decision Support: A Case Study in Sales Forecasting,” in *Proc. IEEE Int. Conf. Artificial Intelligence and Data Analytics (CAIDA)*, 2023.
- [10] Y. Xiong, “Development of an AI-Driven Model for Drug Sales Prediction Using XGBoost and Enhanced Golden Eagle Optimization (EGEO),” *Informatica*, vol. 49, 2025.